Download Latest PDF Resume

Profile Summary

Lead Data Scientist and AI technical lead with 8+ years of experience developing, evaluating and productionizing machine-learning systems across NLP, recommendation, ranking, forecasting, dynamic pricing and generative AI. Combines formal training in data analysis and statistics, advanced machine learning and NLP with end-to-end ownership from ambiguous problem formulation through experimentation, deployment, governance, monitoring and stakeholder adoption. Translates business questions into valid statistical and ML problems, tests assumptions, establishes rigorous baselines and independent evaluations, and determines when evidence is strong enough to support deployment—and when a model output should not be used. Recent work spans enterprise RAG, LLM and agent evaluation, multi-agent orchestration and MCP-based AI platforms, applying modern AI capabilities within a broader discipline of reliable, reproducible and decision-relevant system design.

Education

- Data Analysis / Statistics | Nackademin (*2015 - 2017*)
- International Relations | Stockholms University (*2010 - 2015*)

Advanced Technical Specialisations

- Advanced Machine Learning Specialisation | Higher School of Economics (*2020 - 2022*)
- Natural Language Processing Specialisation (*2023*)

Publications, Projects & Podcasts

1. Building Baseline Models with Design Patterns in ML (Part 1). Publication
2. Exploring Stable Diffusion in Google Colab using CUDA: A Step-by-Step Tutorial. Publication
3. Hyperdimensional Computing: Taking AI to the Next Level by Emulating the Brain. Publication
4. Data Ethics and Responsible AI. Podcast
5. Building a Comprehensive RAG Application: Frontend and Backend Development with Azure OpenAI, Flask & LangChain. Publication

More of my articles can be found on my blog: Link to my blog

Key Skills

- **Statistical Modeling & Experimentation:** Bayesian and probabilistic modeling, A/B testing, experimentation frameworks, time-series modeling and forecasting, model evaluation and

experiment tracking, regression and classification, Markov-chain attribution modeling, feature engineering, exploratory data analysis
- **Programming Languages & Libraries:** Python, SQL, R, Pandas, NumPy, TensorFlow, PyTorch, Scikit-Learn, Lifetimes Python Package, MLFlow, Streamlit
- **NLP and generative AI:** Hugging Face, NLTK, SpaCy, LangChain, LangGraph, AutoGen
- **Machine Learning & Data Science:** Reinforcement Learning, Deep Learning, Transfer Learning, Federated Learning, Recommendation Systems, Model Optimization, Model Tracking, Dynamic Pricing
- **Cloud Services:** Snowflake, AWS, Azure, GCP
- **Development Tools & Version Control:** GIT, GitLab, JIRA, Trello, Asana
- **Visualization Tools:** Tableau, Google Data Studio, Power BI, Looker, D3.js
- **MLOps and Deployment:** CI/CD Pipelines, Kubernetes, Docker
- **Big Data Technologies:** Apache Spark, Hadoop
- **Data Governance and Compliance:** GDPR, CCPA

Soft Skills

- Effective mentor to junior data scientists, interns, and university students, helping them develop skills in Python, ML fundamentals, and production workflows.
- Strong stakeholder communication, translating complex technical concepts into actionable business insights.
- Leadership experience in guiding project direction, reviewing code, and ensuring best practices across teams.
- Adaptability in fast-paced environments with changing priorities and emerging technologies.

Work Experience

Lead Data Scientist – AI Initiatives | Stora Enso (*January 2025 - Present*)

Enterprise Multi-Agent Intelligence Platform (2025 – Present)

- Architected and built a pilot multi-tenant platform, currently in active development, that answers natural-language business questions by selecting registered agents or dynamically generating specialized LLM agents from a governed tool catalog.
- Designed both single-agent execution and multi-agent business-case orchestration through decomposition, assignment, execution, and synthesis.
- Built an MCP server and gateway over Streamable HTTP, with external-connector support, per-connection failure isolation, tool allowlists, and tiered access controls.
- Implemented allowlisted SQL execution, row-level tenant isolation, audit logging, and sampled LLM-judge evaluations.
- Delivered real-time SSE streaming between FastAPI and Next.js/React, exposing live agent and tool-call progress.
- Established a CI-driven agentic SDLC in which maintenance

agents address bugs, dependency upgrades, refactoring, and documentation/test updates through pull requests, backed by approximately 2,400 automated tests.

Stack: Python, FastAPI, PostgreSQL/PGVector, MCP, Azure OpenAI, Docker, Azure Key Vault, Next.js, React

- Provide technical leadership across AI initiatives, from problem definition and architecture through validation and production delivery.
- Mentor data scientists and engineers through architectural guidance, code reviews, and production-oriented development practices.

Senior Data Scientist | Stora Enso (_June 2022 - 2025)

AI Negotiation Agent for Stora Enso (2024 – Present)

- Designed and deployed a production multi-agent system, currently operating in monitoring mode, coordinating 5+ specialized LLM agents through a custom orchestrator extending Magentic-One/AutoGen.
- Implemented source-grounded claim verification and an evaluation framework using LLM judges and RAGAS.
- Built PostgreSQL/PGVector RAG for financial and customer documents, supported by a Snowflake-to-PostgreSQL profitability-analytics pipeline.
- Deployed the system on Azure using Azure OpenAI, Key Vault, and Docker.
- Supported approximately 10 stakeholders, including decision-makers; system outputs have been used in formal negotiations.

Stack: Python, FastAPI, AutoGen, PostgreSQL/PGVector, Snowflake, Azure OpenAI, Docker

AI Acceleration with Generative AI Project (January 2024 - 2025):

- Designed and delivered production RAG solutions supporting internal teams, including a protected use case that enabled employees to interact with legal documents.
- Led end-to-end implementation using LangChain, Hugging Face, Azure OpenAI, Flask, Docker, PostgreSQL, and Azure.
- Researched LLM reasoning and multi-agent systems and translated relevant findings into practical generative-AI applications.

Topic Modelling for Safety notifications(January 2024 - December 2024): - Managed extensive volumes of safety notifications in textual format, utilizing topic modelling techniques to identify themes and generate clusters. - Employed models such as LDA, BERT, K-means, and hybrid variations to extract insights from the data efficiently.

Precision Forestry Project (August 2023 - December 2023):

- Optimized code structure using design patterns.
- Refactored regression models and implemented Azure AutoML

SDK.
- Streamlined development processes, enhancing overall codebase efficiency.

Selfly Store Project (June 2022 - June 2023):

Generated product ideas and developed forecasting, recommendation, and dynamic-pricing models for AI-enabled unattended retail systems. Was recognized internally as an inventor in a company patent initiative; the US patent process is no longer active. Project Stages:

A. Analytics/Engineering Stage:

- Managed and processed extensive cloud-stored data.
- Designed a preprocessing pipeline for product inventory.
- Improved product categorization using a logic-based approach and TF-IDF-based classifier.
- Created a cabinet lifetime value model for performance assessment.

B. Artificial Intelligence Stage:

- Delivered forecasting models to production using LSTM and ensemble methods.
- Developed deep neural network-based models for propensity score.
- Developed collaborative-filtering and hybrid recommender systems.
- Developed Bayesian approaches for dynamic-pricing decisions.

Tech Stack: Databricks, MLFlow, Hugging Face (Open-source NLP library), Lanchain, Streamlit, Flask, Docker Containers, Azure Open AI, PostgreSQL, Azure AutoML SDK, Scikit-Learn, Python, SQL, SnowFlake, GIT, Azure Cloud, MongoDB

Data Scientist | TUI Global (*March 2021 - June 2022*)

Customer Life Time Value Project (Mar 2021 - Nov 2021):

- Developed Customer Lifetime Value models to support business decisions.
- Implemented probabilistic and neural-network approaches using the Lifetimes Python package.
- Stored model outputs in Snowflake and shared results through a Tableau dashboard.

Recommender Systems Project (Mar 2021 - Oct 2021):

- Built item-based collaborative-filtering and user-based probabilistic models to personalize destination and hotel recommendations.
- Used XGBoost for classification and produced recommendations for email delivery.
- Stored model outputs in Snowflake.

Hotel Search Ranking Model Project (Mar 2021 - Jan 2022):

- Developed a hotel-ranking solution combining booking and clickstream signals.
- Built a stacked approach using Random Forest and XGBoost

classifiers and tracked model results in MLflow.

ML Pipeline Automation Package POC (Mar 2021 - June 2022):

- Developed reusable Python functions and classes covering machine-learning pipeline stages.
- Automated recurring workflow steps and stored outputs in MLflow and local folders for analysis.

Tech Stack: Lifetimes Python Package (Probabilistic modeling), Snowflake, Tableau, XGBoost, MLFlow (Model tracking), Python, SQL, Google Cloud, Amazon Web Services, GitLab

Data Science & Data Analytics | TUI Nordics (*April 2019 - March 2021*)

Machine Learning & Data Science:

- Built attribution models using Markov chains for revenue allocation and optimization.
- Conducted customer segmentation via RFM key metrics: recency, frequency, and monetary value.
- Implemented classification and regression algorithms such as KNN, Random Forest, and XGBoost.
- Engaged in error monitoring, feature engineering, and exploratory data analysis.
- Developed recommendation engines and evaluated recommender systems.
- Applied time series modeling techniques including LSTM, XGBoost, and ARIMA variations.

Data/Web Analysis:

- Provided weekly reporting and analysis to stakeholders for website optimization.
- Utilized SQL (BigQuery) and Python for data analysis.
- Extracted specific data from databases using SQL and analyzed further using Python, Jupyter Notebooks, and PyCharm.
- Created data monitoring dashboards using Google Data Studio or Tableau based on task objectives.

Tech Stack: Snowflake, Tableau, Python, SQL, Google Cloud, Amazon Web Services, GitLab, Google Cloud, Google Data Studio

Guest Lecturer (Data Analytics/Pandas Programming) | Hyper Island (*December 2020*)

- Delivered an online guest lecture for data-analytics students on modern analysis workflows with Python and Pandas.

Tech Stack: Python, Pandas, Jupyter Notebook

Data Analyst - E-commerce | Off The Wall (OTW) (*June 2017 - April 2019*)

- Conducted web/marketing analysis, optimizing website performance and content.
- Provided analytical support for business development and editorial teams.
- Implemented A/B testing and conversion optimization strategies.
- Managed Google Adwords campaigns, conducted keyword analysis, and developed reports from Google Analytics.
- Utilized Google Tag Manager for tag management and analyzed SEO performance.
- Applied advanced statistical techniques for predictive analysis.
- Collaborated with social media team, managing ads, optimizing campaigns, and analyzing data using Facebook Analytics and Facebook pixel.

Contact info:

- **Email:** tabers77@gmail.com
- **GitHub:** GitHub Profile
- **Data Science Portfolio:** GitHub Repository
- **LinkedIn:** LinkedIn Connect
- **Digital Resume & Portfolio:** [Digital Resume & Portfolio])