

[Download Latest PDF Resume](#)

Profile Summary

Senior Data Scientist and AI technical lead with 8+ years of experience taking machine-learning and applied-AI initiatives from problem framing and model selection through production delivery. Recent work includes architecting a pilot multi-tenant agent platform, MCP-based tool integration, production multi-agent systems, RAG for sensitive enterprise knowledge, LLM evaluation, forecasting, and recommendation systems. Combines hands-on Python and cloud engineering with technical leadership, mentoring, and clear communication with business and engineering stakeholders.

Education

- Advance Machine Learning Specialisation | Higher School of Economics (2020 - 2022)
- Data Analysis/Statistics | Nackademin (2015 - 2017)
- International Relations | Stockholms University (2010 - 2015)

Courses & Specialisations

- Advance Machine Learning Specialisation | Higher School of Economics (2020 - 2022)
- [Natural Language Processing Specialisation \(2023\)](#)

Publications, Projects & Podcasts

1. Building Baseline Models with Design Patterns in ML (Part 1). [Publication](#)
2. Exploring Stable Diffusion in Google Colab using CUDA: A Step-by-Step Tutorial. [Publication](#)
3. Hyperdimensional Computing: Taking AI to the Next Level by Emulating the Brain. [Publication](#)
4. Data Ethics and Responsible AI. [Podcast](#)
5. Building a Comprehensive RAG Application: Frontend and Backend Development with Azure OpenAI, Flask & LangChain. [Publication](#)

More of my articles can be found on my blog: [Link to my blog](#)

Key Skills

- **Programming Languages & Libraries:** Python, SQL, R, Pandas, NumPy, TensorFlow, PyTorch, Scikit-Learn, Lifetimes Python Package, MLFlow, Streamlit
- **NLP and generative AI:** Hugging Face, NLTK, SpaCy, LangChain, LangGraph, AutoGen
- **Machine Learning & Data Science:** Reinforcement Learning, Deep Learning, Transfer Learning, Federated Learning,

Recommendation Systems, Model Optimization, Model Tracking, Time Series Modeling, Feature Engineering, EDA, Bayesian Statistics, Dynamic Pricing

- **Cloud Services:** Snowflake, AWS, Azure, GCP
- **Development Tools & Version Control:** GIT, GitLab, JIRA, Trello, Asana
- **Visualization Tools:** Tableau, Google Data Studio, Power BI, Looker, D3.js
- **MLOps and Deployment:** CI/CD Pipelines, Kubernetes, Docker
- **Big Data Technologies:** Apache Spark, Hadoop
- **Data Governance and Compliance:** GDPR, CCPA
- **Experimentation and Testing:** A/B Testing, Experimentation Frameworks

Soft Skills

- Effective mentor to junior data scientists, interns, and university students, helping them develop skills in Python, ML fundamentals, and production workflows.
- Strong stakeholder communication, translating complex technical concepts into actionable business insights.
- Leadership experience in guiding project direction, reviewing code, and ensuring best practices across teams.
- Adaptability in fast-paced environments with changing priorities and emerging technologies.

Work Experience

Lead Data Scientist - AI Initiatives | Stora Enso (*January 2025 - Present*)

Enterprise Multi-Agent Intelligence Platform (2025 - Present)

- Architected and built a pilot multi-tenant platform, currently in active development, that answers natural-language business questions by selecting registered agents or dynamically generating specialized LLM agents from a governed tool catalog.
- Designed both single-agent execution and multi-agent business-case orchestration through decomposition, assignment, execution, and synthesis.
- Built an MCP server and gateway over Streamable HTTP, with external-connector support, per-connection failure isolation, tool allowlists, and tiered access controls.
- Implemented allowlisted SQL execution, row-level tenant isolation, audit logging, and sampled LLM-judge evaluations.
- Delivered real-time SSE streaming between FastAPI and Next.js/React, exposing live agent and tool-call progress.
- Established a CI-driven agentic SDLC in which maintenance agents address bugs, dependency upgrades, refactoring, and documentation/test updates through pull requests, backed by approximately 2,400 automated tests.

Stack: Python, FastAPI, PostgreSQL/PGVector, MCP, Azure OpenAI, Docker, Azure Key Vault, Next.js, React

- Provide technical leadership across AI initiatives, from problem definition and architecture through validation and production delivery.
- Mentor data scientists and engineers through architectural guidance, code reviews, and production-oriented development practices.

Senior Data Scientist | Stora Enso (June 2022 - 2025)

AI Negotiation Agent for Stora Enso (2024 - Present)

- Designed and deployed a production multi-agent system, currently operating in monitoring mode, coordinating 5+ specialized LLM agents through a custom orchestrator extending Magentic-One/AutoGen.
- Implemented source-grounded claim verification and an evaluation framework using LLM judges and RAGAS.
- Built PostgreSQL/PGVector RAG for financial and customer documents, supported by a Snowflake-to-PostgreSQL profitability-analytics pipeline.
- Deployed the system on Azure using Azure OpenAI, Key Vault, and Docker.
- Supported approximately 10 stakeholders, including decision-makers; system outputs have been used in formal negotiations.

Stack: Python, FastAPI, AutoGen, PostgreSQL/PGVector, Snowflake, Azure OpenAI, Docker

AI Acceleration with Generative AI Project (January 2024 - 2025):

- Designed and delivered production RAG solutions supporting internal teams, including a protected use case that enabled employees to interact with legal documents.
- Led end-to-end implementation using LangChain, Hugging Face, Azure OpenAI, Flask, Docker, PostgreSQL, and Azure.
- Researched LLM reasoning and multi-agent systems and translated relevant findings into practical generative-AI applications.

Topic Modelling for Safety notifications(January 2024 - December 2024): - Managed extensive volumes of safety notifications in textual format, utilizing topic modelling techniques to identify themes and generate clusters. - Employed models such as LDA, BERT, K-means, and hybrid variations to extract insights from the data efficiently.

Precision Forestry Project (August 2023 - December 2023):

- Optimized code structure using design patterns.
- Refactored regression models and implemented Azure AutoML SDK.
- Streamlined development processes, enhancing overall codebase efficiency.

Selfly Store Project (June 2022 - June 2023):

Generated product ideas and developed forecasting, recommendation, and dynamic-pricing models for AI-enabled unattended retail systems. Was recognized internally as an inventor in a company patent initiative; the US patent process is no longer active. Project Stages:

A. Analytics/Engineering Stage:

- Managed and processed extensive cloud-stored data.
- Designed a preprocessing pipeline for product inventory.
- Improved product categorization using a logic-based approach and TF-IDF-based classifier.
- Created a cabinet lifetime value model for performance assessment.

B. Artificial Intelligence Stage:

- Delivered forecasting models to production using LSTM and ensemble methods.
- Developed deep neural network-based models for propensity score.
- Developed collaborative-filtering and hybrid recommender systems.
- Developed Bayesian approaches for dynamic-pricing decisions.

Tech Stack: Databricks, MLFlow, Hugging Face (Open-source NLP library), Lanchain, Streamlit, Flask, Docker Containers, Azure Open AI, PostgreSQL, Azure AutoML SDK, Scikit-Learn, Python, SQL, Snowflake, GIT, Azure Cloud, MongoDB

Data Scientist | TUI Global (*March 2021 - June 2022*)

Customer Life Time Value Project (Mar 2021 - Nov 2021):

- Developed Customer Lifetime Value models to support business decisions.
- Implemented probabilistic and neural-network approaches using the Lifetimes Python package.
- Stored model outputs in Snowflake and shared results through a Tableau dashboard.

Recommender Systems Project (Mar 2021 - Oct 2021):

- Built item-based collaborative-filtering and user-based probabilistic models to personalize destination and hotel recommendations.
- Used XGBoost for classification and produced recommendations for email delivery.
- Stored model outputs in Snowflake.

Hotel Search Ranking Model Project (Mar 2021 - Jan 2022):

- Developed a hotel-ranking solution combining booking and clickstream signals.
- Built a stacked approach using Random Forest and XGBoost classifiers and tracked model results in MLflow.

ML Pipeline Automation Package POC (Mar 2021 - June 2022):

- Developed reusable Python functions and classes covering machine-learning pipeline stages.

- Automated recurring workflow steps and stored outputs in MLflow and local folders for analysis.

Tech Stack: Lifetimes Python Package (Probabilistic modeling), Snowflake, Tableau, XGBoost, MLFlow (Model tracking), Python, SQL, Google Cloud, Amazon Web Services, GitLab

Data Science & Data Analytics | TUI Nordics (*April 2019 - March 2021*)

Machine Learning & Data Science:

- Built attribution models using Markov chains for revenue allocation and optimization.
- Conducted customer segmentation via RFM key metrics: recency, frequency, and monetary value.
- Implemented classification and regression algorithms such as KNN, Random Forest, and XGBoost.
- Engaged in error monitoring, feature engineering, and exploratory data analysis.
- Developed recommendation engines and evaluated recommender systems.
- Applied time series modeling techniques including LSTM, XGBoost, and ARIMA variations.

Data/Web Analysis:

- Provided weekly reporting and analysis to stakeholders for website optimization.
- Utilized SQL (BigQuery) and Python for data analysis.
- Extracted specific data from databases using SQL and analyzed further using Python, Jupyter Notebooks, and PyCharm.
- Created data monitoring dashboards using Google Data Studio or Tableau based on task objectives.

Tech Stack: Snowflake, Tableau, Python, SQL, Google Cloud, Amazon Web Services, GitLab, Google Cloud, Google Data Studio

Guest Lecturer (Data Analytics/Pandas Programming) | Hyper Island (*December 2020*)

- Delivered an online guest lecture for data-analytics students on modern analysis workflows with Python and Pandas.

Tech Stack: Python, Pandas, Jupyter Notebook

Data Analyst - E-commerce | Off The Wall (OTW) (*June 2017 - April 2019*)

- Conducted web/marketing analysis, optimizing website performance and content.
- Provided analytical support for business development and editorial teams.
- Implemented A/B testing and conversion optimization strategies.

- Managed Google Adwords campaigns, conducted keyword analysis, and developed reports from Google Analytics.
- Utilized Google Tag Manager for tag management and analyzed SEO performance.
- Applied advanced statistical techniques for predictive analysis.
- Collaborated with social media team, managing ads, optimizing campaigns, and analyzing data using Facebook Analytics and Facebook pixel.

Contact info:

- **Email:** tabers77@gmail.com
- **GitHub:** [GitHub Profile](#)
- **Data Science Portfolio:** [GitHub Repository](#)
- **LinkedIn:** [LinkedIn Connect](#)
- **Digital Resume & Portfolio:** [\[Digital Resume & Portfolio\]](#)